

Superforecaster — Kalshi Soft Markets

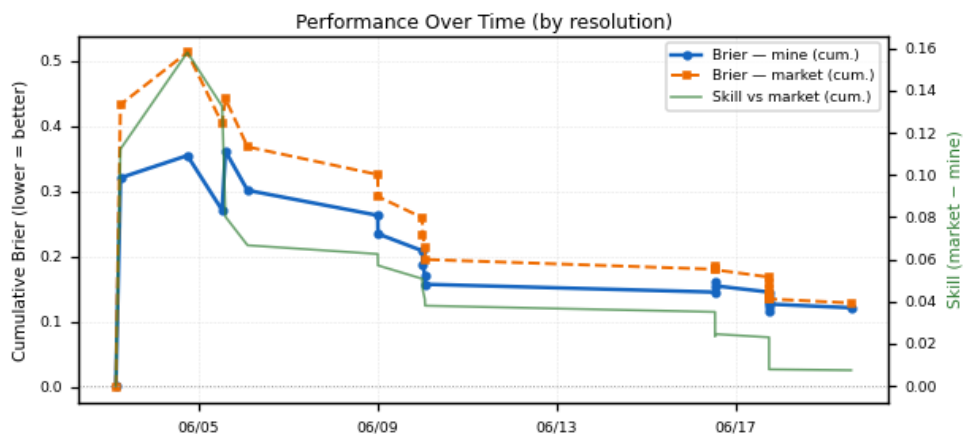
Generated: 2026-06-19T19:16:08Z

Metric	Value
Active markets	17
Resolved markets	22
Total forecast entries	90
Profitable leans (after fees)	1
Resolved with calibration	22
Brier (mine, mean)	0.1209
Brier (market, mean)	0.1283
Skill vs market	0.0074

Performance Over Time

Resolved: **22** Brier (mine): **0.121** vs market **0.128** Skill: **+0.007** (beating the market)

Provisional — only 22 resolved market(s). A skill gap is not statistically distinguishable from luck below ~30 resolutions; read the trend, not the point estimate.



Where the edge is — by segment

Strongest: **politics / us-governor-primary** (skill +0.074, n=3). Weakest: **economy / fed-rates** (skill -0.034, n=6).

By category

Category	n	Brier (mine)	Brier (mkt)	Skill
politics	15	0.153	0.174	+0.021
economy	7	0.051	0.030	-0.021

By sub-category

Category / sub-category	n	Brier (mine)	Brier (mkt)	Skill
politics / foreign-election	1	0.423	0.672	+0.250
politics / us-governor-primary	3	0.214	0.289	+0.074
economy / jobs-unemployment	1	0.020	0.076	+0.056

politics / mayoral	2	0.034	0.063	+0.028
politics / us-governor	1	0.000	0.000	+0.000
politics / legislation	1	0.002	0.000	-0.002
politics / foreign-policy	3	0.145	0.116	-0.029
politics / us-senate	4	0.183	0.150	-0.033
economy / fed-rates	6	0.057	0.023	-0.034

Profit & Loss (realized)

Resolved trades: **5** Realized P&L: **-\$0.60** per contract ROI: **-16.7%** Win rate: **60%**

Paper trade: buy the lean side at its entry ask, after Kalshi fees. A high win rate with negative ROI means winners are underpriced relative to the losers — calibration without profit.

By category

Category	Trades	P&L	ROI	Win%	MaxDD
economy	1	+\$0.25	+33.3%	100%	+\$0.00
politics	4	-\$0.85	-29.8%	50%	-\$0.85

Strategy Scoreboard

Every forecast is tagged with the strategy arm that produced it; resolutions score that arm on both Brier skill and realized profit. The topology is measured, not assumed.

Strategy	n	Brier	Skill	Trades	ROI
untagged	13	0.145	+0.035	4	-2.6%
L0-local-single	2	0.031	+0.020	0	—
S0-single	5	0.001	-0.000	0	—
S2-ensemble3-crowd	2	0.354	-0.165	1	-100.0%

Profitability — when does the edge pay? (counterfactual)

Every resolved forecast scored as a modal-side paper trade (22 total, vs only the taken leans in the realized section above). Overall counterfactual ROI: **-7.4%**, win rate 82%. The breakdowns below isolate **when** the strategy is profitable — the goal is profit, so this is the scoreboard that matters most.

Confidence gate: taken leans vs the ones we declined

Decision	n	CF P&L	CF ROI	Win%	MaxDD
taken	5	+\$0.32	+8.7%	80%	-\$0.80
gated/none	17	-\$1.76	-11.1%	82%	-\$1.79

By edge vs market (|my prob – market|) — where profit survives fees

Edge band	n	CF P&L	CF ROI	Win%	MaxDD
0-5pt	15	+\$0.76	+5.3%	100%	-\$0.01
5-10pt	3	-\$0.27	-11.9%	67%	-\$0.80
10-20pt	3	-\$1.53	-60.5%	33%	-\$1.78
20pt+	1	-\$0.39	-100.0%	0%	-\$0.39

By stated confidence — does confidence predict profit?

Confidence	n	CF P&L	CF ROI	Win%	MaxDD
high	14	+\$0.38	+2.8%	100%	-\$0.01
medium	5	-\$0.87	-22.5%	60%	-\$1.74
low	3	-\$0.94	-48.5%	33%	-\$0.94

By category

Category	n	CF P&L	CF ROI	Win%	MaxDD
economy	7	-\$0.11	-1.8%	86%	-\$0.39
politics	15	-\$1.33	-10.0%	80%	-\$2.59

By strategy arm

Strategy	n	CF P&L	CF ROI	Win%	MaxDD
L0-local-single	2	+\$0.29	+17.0%	100%	+\$0.00
S2-ensemble3-crowd	2	+\$0.03	+3.1%	50%	-\$0.39
S0-single	5	+\$0.03	+0.6%	100%	+\$0.00
untagged	13	-\$1.79	-15.2%	77%	-\$2.34

Open positions — expected P&L (forward-looking)

Forward-looking expected P&L on **52** open modal-side positions (mark-to-model): **-\$0.22/contract** total (19 are +EV). Of that, **+\$0.84** is on actionable leans (the rest is counterfactual — positions the confidence gate is holding us out of).

Provisional — 22 resolved. Profit ROI is even noisier than Brier at small n; read the direction, not the digits.

Autonomous Learning — policy & proposals

The decision policy (when to take a position) is data, not frozen code. The learning pass reads the system's own resolved record and proposes nudges, but anti-overfit guardrails (min-n, max-step) gate every change. **Policy v1**; proposals as of 2026-06-19.

Knob (decision threshold)	Value
min profitable EV (\$/contract)	0.020
max market-fade gap (no-high-conf)	0.200
conviction medium / high EV	0.05 / 0.12
low confidence never leans	True
adversarial veto binding	True

Pending proposals (from the system's own outcomes)

Knob	Current → Proposed	n	Guardrail	Rationale
max_market_disagreement	0.2 → 0.1	4	INSUFFICIENT_DATA	Counterfactual ROI negative in 10-20pt, 20pt+ -> wide fades lose; tighten
low_confidence_never_leans	True → True	19	INSUFFICIENT_DATA	medium-confidence trades lose while high-confidence profit -> consider g
adversarial_veto_binding	True → None	1	INSUFFICIENT_DATA	No resolved positions carry an adversarial verdict yet — the gate is new.

INSUFFICIENT_DATA = directionally real but too few samples to trust; HUMAN_GATE = change too large to auto-apply; AUTO_OK = guardrail-cleared, auto-applied. Nothing auto-applies until the record is large enough.

Active Markets

Will Toby Doeden be the Republican nominee for Governor in South Dakota?

(KXGOVSDNOMR-26-TDOE)

Category: politics | Closes: 2026-06-30 14:00 UTC

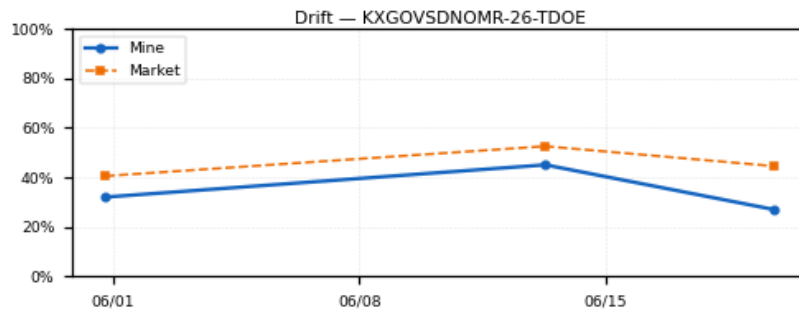
My prob: **27.0%** Confidence: **medium** Market: **44.5%** Edge: **-17.5%**

Lean: **NO** Conviction: **high**

Trade (spot): BUY NO @ \$0.56 now (taker, fee \$0.02/contract) -> net EV **+\$0.15/contract**

Rationale: SD primary Jun 2 hit no 35% threshold; 2-way Jul 28 runoff Doeden (31% plurality) vs incumbent Gov Rhoden (25%). ~69% did not back Doeden; anti-Doeden (Johnson) base expected to consolidate behind incumbent. Doeden a clear underdog with intense base.

Key drivers: Jul 28 runoff vs incumbent Rhoden; 69% did not back Doeden; Johnson voters consolidate to Rhoden; Doeden base intensity in low-turnout runoff



Will the 7-day moving average of transit calls through the Strait of Hormuz as reported by the IMF PortWatch be above 60 before July 1, 2026?

(KXHORMUZNORM-26MAR17-B260701)

Category: politics | Closes: 2026-06-15 14:00 UTC


My prob: **1.2%** Confidence: **high** Market: **7.3%** Edge: **-6.1%**

Lean: **NONE** Conviction: **low**

No actionable lean: ADVERSARIAL VETO: The forecaster's 'NO' position assumes a future crisis with no present indicators, diverging sharply from the market's 7.3% implied probability. The 5-ships vs 60 threshold gap is extreme but uncorroborated by liquid market data. Overconfid (Indicative raw EV +\$0.05/contract at my probability, not recommended.)

Rationale: 2026 Strait of Hormuz crisis collapsed transit >90%; IMF PortWatch 7-day MA ~5 ships vs the 60 threshold, no recovery before the window. Effectively determined NO.

Key drivers: Hormuz crisis collapsed traffic >90%; 7-day MA ~5 vs 60 threshold; no recovery before window; US-Iran deal too late



Will Donald J. Trump attend The 2026 FIFA World Cup Final? (KXTRUMPATTEND)

Category: politics | Closes: 2026-07-20 14:00 UTC


My prob: **85.0%** Confidence: **medium** Market: **81.5%** Edge: **3.5%**

Lean: **NONE** Conviction: **low**

No actionable lean: I lean YES (my prob 85%) but it isn't +EV at the ask — no value bet (betting the other side would contradict my forecast). (Indicative raw EV +\$0.01/contract at my probability, not recommended.)

Rationale: Trump plans to attend (WH 'holding' his schedule); FIFA wants him to present the trophy; precedent at Club World Cup; marquee US-hosted event he is motivated to attend. Residual scheduling/security/health risk a month out.

Key drivers: WH holding the date; FIFA trophy-presentation role; Club World Cup precedent; strong personal motivation



Will Marco Rubio visit Iran before Jul 1, 2026? (KXVISITIRAN-26JUL01-MRUB)

Category: politics | Closes: 2026-07-01 14:00 UTC


My prob: **1.0%** Confidence: **high** Market: **0.9%** Edge: **0.1%**

Lean: **NONE** Conviction: **low**

No actionable lean: I lean NO (my prob 1%) but it isn't +EV at the ask — no value bet (betting the other side would contradict my forecast).

Rationale: No US Secretary of State has visited Iran since 1979; 2026 talks held in neutral venues (Geneva); no plan to travel to Tehran.

Key drivers: no SoS visit since 1979; neutral-venue talks; no Tehran plan



Will JD Vance visit Iran before Jul 1, 2026? (KXVISITIRAN-26JUL01-JVAN)

Category: politics | Closes: 2026-07-01 14:00 UTC


My prob: **2.0%** Confidence: **medium** Market: **1.1%** Edge: **0.9%**

Lean: **NONE** Conviction: **low**

No actionable lean: I lean NO (my prob 2%) but it isn't +EV at the ask — no value bet (betting the other side would contradict my forecast).

Rationale: Vance's stated plan is Switzerland, not Iran; no US VP has visited Iran since 1979; talks in neutral venues.

Key drivers: Vance plan is Switzerland not Iran; no VP visit since 1979; neutral venues



Will the unemployment rate (U-3) be above 4.1% in June? (KXU3-26JUN-T4.1)

Category: economy | Closes: 2026-07-02 14:00 UTC


My prob: **92.0%** Confidence: **high** Market: **93.0%** Edge: **-1.0%**

Lean: **NONE** Conviction: **low**

No actionable lean: I lean YES (my prob 92%) but it isn't +EV at the ask — no value bet (betting the other side would contradict my forecast).

Rationale: May 2026 U-3 = 4.3%, in a 4.3-4.5% range since Jul 2025; no sign of a drop below 4.1% in June. Very likely above 4.1%.

Key drivers: May 4.3%; stable 4.3-4.5% range; no rapid improvement



Will Trump make 2 trips to Mar-a-Lago as President in Jun 2026? (KXLAGODAYS-26JUN-2)

Category: politics | Closes: 2026-06-30 14:00 UTC


My prob: **3.0%** Confidence: **medium** Market: **0.5%** Edge: **2.5%**

Lean: **NONE** Conviction: **low**

No actionable lean: I lean NO (my prob 3%) but it isn't +EV at the ask — no value bet (betting the other side would contradict my forecast).

Rationale: Two June Mar-a-Lago trips highly unlikely given summer Bedminster pattern + Iran war in DC; corrected down from incoherent local estimate.

Key drivers: summer Bedminster; Iran war; 2 trips needs unlikely double travel



Will any U.S. House member visit Iran before Jul 1, 2026? (KXVISITIRAN-26JUL01-AHOU)

Category: politics | Closes: 2026-07-01 14:00 UTC


My prob: **3.0%** Confidence: **medium** Market: **0.7%** Edge: **2.3%**

Lean: **NONE** Conviction: **low**

No actionable lean: I lean NO (my prob 3%) but it isn't +EV at the ask — no value bet (betting the other side would contradict my forecast).

Rationale: All US-Iran engagement is in neutral venues (Geneva/Switzerland/Rome), not Tehran; no House member traveling to Iran; physically near-unprecedented. Corrected down from local overestimate.

Key drivers: talks in neutral venues not Tehran; no House travel to Iran; near-unprecedented



Will the US agree to a new Iranian nuclear deal this year? (KXUSAIRANAGREEMENT-27-26JUL)

Category: politics | Closes: 2026-07-01 14:00 UTC


My prob: **5.0%** Confidence: **medium** Market: **0.5%** Edge: **4.5%**

Lean: **NONE** Conviction: **low**

No actionable lean: I lean NO (my prob 5%) but it isn't +EV at the ask — no value bet (betting the other side would contradict my forecast).

Rationale: June 14-19 MOU is explicitly a roadmap, not a nuclear deal; opens a 60-day nuclear negotiation window ending ~mid-August — after the July 1 close. Talks contentious (missiles, rebuilding).

Key drivers: MOU is roadmap not deal; 60-day window ends mid-Aug past close; contentious nuclear talks



Will legislation that reauthorizes FISA Section 702 authority become law before Jun 25, 2026? (KXFISAEXTEND-26JUN-26JUN25)

Category: politics | Closes: 2026-07-01 14:00 UTC


My prob: **6.0%** Confidence: **medium** Market: **1.5%** Edge: **4.5%**

Lean: **NONE** Conviction: **low**

No actionable lean: I lean NO (my prob 6%) but it isn't +EV at the ask — no value bet (betting the other side would contradict my forecast).

Rationale: 702 in active lapse: House failed short extension 198-218, Senate stalled 47-52, 'lapse assured', Dems blocking over Pulte DNI. No path to law in 6 days.

Key drivers: House extension failed 198-218; Senate motion failed 47-52; program lapsed; Pulte standoff



Will Trump make 1 trips to Mar-a-Lago as President in Jun 2026? (KXLAGODAYS-26JUN-1)

Category: politics | Closes: 2026-06-30 14:00 UTC


My prob: **13.0%** Confidence: **medium** Market: **5.5%** Edge: **7.5%**

Lean: **NONE** Conviction: **low**

No actionable lean: I lean NO (my prob 13%) but it isn't +EV at the ask — no value bet (betting the other side would contradict my forecast).

Rationale: Coherence with P(0 trips)~0.82: residual mass for exactly 1 trip is modest. Summer/war suppress Florida travel.

Key drivers: coherence with high P(0); summer suppression of FL travel



Will the President sign more than 0 Executive Orders between Jun 7, 2026 and Jun 13, 2026? (KXEOWEEK-26JUN13-0)

Category: politics | Closes: 2026-06-27 14:00 UTC


My prob: **10.0%** Confidence: **medium** Market: **0.5%** Edge: **9.5%**

Lean: **NONE** Conviction: **low**

No actionable lean: I lean NO (my prob 10%) but it isn't +EV at the ask — no value bet (betting the other side would contradict my forecast).

Rationale: Window Jun 7-13 is past. Numbered EO series stopped at EO 14411 on Jun 3 (pre-window); in-window actions were a proclamation (Jun 11) + memorandum (Jun 12), not numbered EOs -> strict reading NO. 0.10 reflects residual Federal-Register-lag and resolution-ambiguity tail risk vs market 0.5%.

Key drivers: numbered EO series stopped at 14411 Jun 3; in-window actions were proclamation+memo not EOs; Federal Register lag risk; resolution-definition ambiguity



Will Trump make 0 trips to Mar-a-Lago as President in Jun 2026? (KXLAGODAYS-26JUN-0)

Category: politics | Closes: 2026-06-30 14:00 UTC


My prob: **82.0%** Confidence: **medium** Market: **95.5%** Edge: **-13.5%**

Lean: **NONE** Conviction: **low**

No actionable lean: I lean YES (my prob 82%) but it isn't +EV at the ask — no value bet (betting the other side would contradict my forecast).

Rationale: Summer = Bedminster not Mar-a-Lago (skips Florida heat); Iran war kept him in DC through the June 14 ceasefire; no reported June trip. 0 trips most likely.

Key drivers: summer Bedminster preference; Iran war kept him in DC; no reported June trip



CPI year-over-year in Jun 2026? (KXECONSTATCPIYOY-26JUN-T4.1)

Category: economy | Closes: 2026-07-14 14:00 UTC


My prob: **20.0%** Confidence: **low** Market: **2.0%** Edge: **18.0%**

Lean: **NONE** Conviction: **low**

No actionable lean: I lean NO (my prob 20%) but it isn't +EV at the ask — no value bet (betting the other side would contradict my forecast).

Rationale: May YoY 4.2%, only 0.1pp above the 4.1% threshold; post-ceasefire energy reversal makes a sub-4.1% June print quite achievable. Local model (0.65) under-weighted that; the two CPI markets coherently price a sharp June disinflation, so I defer substantially while keeping residual chance energy doesn't fully reverse within June.

Key drivers: May YoY 4.2% only 0.1pp over threshold; energy reversal; coherent disinflation signal across both CPI markets



Will Mike Mazzei be the Republican nominee for Governor in Oklahoma?

(KXGOVOKNOMR-26-MMAZ)

Category: politics | Closes: 2026-06-30 14:00 UTC


My prob: 55.0% Confidence: medium Market: 75.0% Edge: -20.0%

Lean: NONE Conviction: low

No actionable lean: I lean YES (my prob 55%) but it isn't +EV at the ask — no value bet (betting the other side would contradict my forecast).

Rationale: OK primary already held Jun 16; now a 2-way Aug 25 runoff Drummond (26.3%) vs Mazzei (26.0%). Mazzei has Trump endorsement + favorable transfer of Merrick's America-First bloc; Drummond led R1 with higher statewide name ID. Net slight Mazzei edge.

Key drivers: Aug 25 runoff after near-tie R1; Trump endorsement of Mazzei; Merrick bloc transfers to Mazzei; Drummond incumbent-AG name ID



Will CPI rise more than 0.0% in June 2026? (KXCPI-26JUN-T0.0)

Category: economy | Closes: 2026-07-14 13:56 UTC


My prob: 40.0% Confidence: low Market: 12.5% Edge: 27.5%

Lean: NONE Conviction: low


No actionable lean: I lean NO (my prob 40%) but it isn't +EV at the ask — no value bet (betting the other side would contradict my forecast).

Rationale: Local model (0.75) under-weighted the post-ceasefire energy reversal: Hormuz reopened June 14, oil/gas falling, a major drag on June headline MoM after May's gas +40.5% spike. Liquid market prices ~87% chance of flat/negative MoM; deferring substantially toward that mechanism-backed signal while retaining upside from sticky core/shelter.

Key drivers: post-ceasefire oil/gas reversal drags headline; May gas spike unwinds; sticky positive core offsets; strong liquid-market disinflation signal



Edge Overview




Profitable Leans (after fees)

Ticker	Side	Spot EV/contract	Limit EV/contract	Conviction
KXGOVSDNOMR-26-TDOENO		+\$0.15	—	high

Fees use Kalshi's standard taker formula $\text{ceil}(0.07 * p * (1 - p)) / \text{contract}$; resting (maker) limit fills may be cheaper on some series. Paper only - read-only key places no orders.

Calibration



Ticker	Title	Outcome	Brier (mine)	Brier (mkt)
KXGOVCA-26-CBIA	Who will win the governorship in Ca...	NO	0.0002	0.0003
KXGOVCAPRIMARY-26-CBIA	Who will win California's top-two p...	NO	0.0016	0.0009
KXGOVCAPRIMARYPARTY-26-CBIA	Will all advance from California's ...	NO	0.0009	0.0000
KXGOVIANOMR-26-RFEE	Will Randy Feenstra be the Republic...	NO	0.6400	0.8649
KXLAMAYOR1R-26-SPRA	Who will win the first round of the...	NO	0.0361	0.0552
KXMAYORLA-26-SPRA	Who will win Los Angeles Mayoral El...	NO	0.0324	0.0702
KXSENATEMED-26-GRA	Will Graham Platner be the Democrat...	YES	0.0025	0.0063
KXSENATENJR-26-VBRA	Will Vinnie Brand be the Republican...	NO	0.0004	0.0000
KXSENATEREC-26MAY-JUN06	Will a reconciliation bill passed t...	YES	0.7225	0.5929
KXSENATESCR-26-LGRA	Will Lindsey Graham be the Republic...	YES	0.0049	0.0009
KXSEOULMAYOR-26JUN03	Will Moon Won-o win the 2026 Seoul...	NO	0.4225	0.6724
KXU3-26MAY-T4.2	Will the unemployment rate (U-3) be...	YES	0.0196	0.0756
KXFEDDECISION-26JUN-0251	Will the Federal Reserve Cut rates ...	NO	0.0009	0.0006
KXFEDDECISION-26JUN-0261	Will the Federal Reserve Cut rates ...	NO	0.0004	0.0000
KXFEDDECISION-26JUN-H011	Will the Federal Reserve Hike rates...	YES	0.0009	0.0012
KXFEDDECISION-26JUN-H251	Will the Federal Reserve Hike rates...	NO	0.0009	0.0002
KXFEDDECISION-26JUN-H261	Will the Federal Reserve Hike rates...	NO	0.0001	0.0000
KXFISAEXTEND-26JUN-26JUN19	Will legislation that reauthorizes ...	NO	0.0016	0.0000
KXFOMCDISSENTCOUNT-26JUN19	Will any dissenting votes at the ne...	YES	0.3364	0.1332
KXHORMUZNORM-26MAR-26JUN15	Will the 60-day moving average of tr...	NO	0.0004	0.0000
KXHORMUZWEEKLY-26JUN15	Will there be more than 25 transit ...	NO	0.3721	0.2450
KXHORMUZWEEKLY-26JUN19	Will there be more than 30 transit ...	NO	0.0625	0.1024

Lessons Learned

When diverging below a liquid market that is pricing a strong institutional base rate, require a specific, verifiable reason; narrative of 'division' is not evidence of a formally recorded dissent. Anchor on the base rate.

2026-06-19 | resolution | KXFOMCDISSENTCOUNT-26JUN-0 | economy [unjustified-market-divergence]
Brier (mine): 0.3364 | Brier (market): 0.1332 | lost to market

During an active supply-disruption crisis, anchor flow estimates to the SUPPRESSED regime, not an assumed recovery, until data confirms normalization; and do not enter a position when divergence from a liquid market is small and rests on an unverified assumption.

2026-06-19 | resolution | KXHORMUZWEEKLY-26JUN14-T25 | politics [unjustified-market-divergence]
Brier (mine): 0.3721 | Brier (market): 0.2450 | lost to market

Never move a forecast materially on one noisy data point (esp. a single poll); weight by reliability and require corroboration. Don't recommend fading a liquid market on a low-confidence estimate.


2026-05-31 | user_feedback | politics [overreaction-and-confidence]

Cost & Usage

Latest Run

Field	Value
Run ID	2026-06-14T03:39:30Z
Status	ok
Markets researched	10
Web searches	2
Web fetches	0
Tool calls	12
Duration (s)	0.0
Est. tokens	—

Usage Trend



Cost proxies; [claude.ai/settings/usage](#) is authoritative for tokens/\$.